



Replace Cisco Nexus 3132Q-V cluster switches

Cluster and storage switches

NetApp
October 14, 2022

This PDF was generated from <https://docs.netapp.com/us-en/ontap-systems-switches/switch-cisco-3132q-v/task-how-to-replace-cisco-nexus-3132q-v-cluster-switches.html> on October 14, 2022. Always check docs.netapp.com for the latest.

Table of Contents

| | |
|---|---|
| Replace Cisco Nexus 3132Q-V cluster switches | 1 |
| How to replace Cisco Nexus 3132Q-V cluster switches | 2 |

Replace Cisco Nexus 3132Q-V cluster switches

You must be aware of certain configuration information, port connections and cabling requirements when you replace Cisco Nexus 3132Q-V cluster switches.

- The Cisco Nexus 3132Q-V cluster switch is supported.
- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the [Cisco® Cluster Network Switch Reference Configuration File Download](#) page.
- The cluster switches use the Inter-Switch Link (ISL) ports e1/31-32.
- The [Hardware Universe](#) contains information about supported cabling to Nexus 3132Q-V switches:
 - The nodes with 10 GbE cluster connections require QSFP optical modules with breakout fiber cables or QSFP to SFP+ copper break-out cables.
 - The nodes with 40/100 GbE cluster connections require supported QSFP/QSFP28 optical modules with fiber cables or QSFP/QSFP28 copper direct-attach cables.
 - The cluster switches use the appropriate ISL cabling: 2x QSFP28 fiber or copper direct-attach cables.
- On Nexus 3132Q-V, you can operate QSFP ports as either 40/100 Gb Ethernet or 4 x10 Gb Ethernet modes.

By default, there are 32 ports in the 40/100 Gb Ethernet mode. These 40 Gb Ethernet ports are numbered in a 2-tuple naming convention. For example, the second 40 Gb Ethernet port is numbered as 1/2. The process of changing the configuration from 40 Gb Ethernet to 10 Gb Ethernet is called *breakout* and the process of changing the configuration from 10 Gb Ethernet to 40 Gb Ethernet is called *breakin*. When you break out a 40/100 Gb Ethernet port into 10 Gb Ethernet ports, the resulting ports are numbered using a 3-tuple naming convention. For example, the breakout ports of the second 40/100 Gb Ethernet port are numbered as 1/2/1, 1/2/2, 1/2/3, 1/2/4.

- On the left side of Nexus 3132Q-V is a set of four SFP+ ports multiplexed to the first QSFP port.

By default, the RCF is structured to use the first QSFP port.

You can make four SFP+ ports active instead of a QSFP port for Nexus 3132Q-V by using the hardware profile `front portmode sfp-plus` command. Similarly, you can reset Nexus 3132Q-V to use a QSFP port instead of four SFP+ ports by using the hardware profile `front portmode qsfp` command.

- You must have configured some of the ports on Nexus 3132Q-V to run at 10 GbE or 40/100 GbE.

You can break-out the first six ports into 4x10 GbE mode by using the `interface breakout module 1 port 1-6 map 10g-4x` command. Similarly, you can regroup the first six QSFP+ ports from breakout configuration by using the `no interface breakout module 1 port 1-6 map 10g-4x` command.

- You must have done the planning, migration, and read the required documentation on 10 GbE and 40/100 GbE connectivity from nodes to Nexus 3132Q-V cluster switches.

The [Cisco Ethernet Switches](#) page has information about the ONTAP and NX-OS versions supported in this procedure.

How to replace Cisco Nexus 3132Q-V cluster switches

Replacing a defective Cisco Nexus 3132Q-V switch in a cluster network is a nondisruptive procedure (NDO), and you must perform a specific sequence of tasks.

Before you begin

- The existing cluster and network configuration must have:
 - The Nexus 3132Q-V cluster infrastructure must be redundant and fully functional on both switches.

The [Cisco Ethernet Switch](#) page has the latest RCF and NX-OS versions on your switches.
 - All cluster ports must be in the `up` state.
 - Management connectivity must exist on both switches.
 - All cluster logical interfaces (LIFs) must be in the `up` state and must not have been migrated.
- The Nexus 3132Q-V replacement switch:
 - Management network connectivity on the replacement switch must be functional.
 - Console access to the replacement switch must be in place.
 - The desired RCF and NX-OS operating system image switch must be loaded onto the switch.
 - Initial customization of the switch must be complete.

About this task

This procedure replaces the second Nexus 3132Q-V cluster switch CL2 with new 3132Q-V switch C2. The examples in this procedure use the following switch and node nomenclature:

- `n1_clus1` is the first cluster logical interface (LIF) connected to cluster switch C1 for node n1.
- `n1_clus2` is the first cluster LIF connected to cluster switch CL2 or C2, for node n1.
- `n1_clus3` is the second LIF connected to cluster switch C2, for node n1.
- `n1_clus4` is the second LIF connected to cluster switch CL1, for node n1.
- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the [Cisco® Cluster Network Switch Reference Configuration File Download](#) page.
- The nodes are n1, n2, n3, and n4. - The examples in this procedure use four nodes: Two nodes use four 10 GB cluster interconnect ports: `e0a`, `e0b`, `e0c`, and `e0d`. The other two nodes use two 40 GB cluster interconnect ports: `e4a` and `e4e`. See the [Hardware Universe](#) for the actual cluster ports on your platforms.

This procedure covers the following scenario:

- The cluster starts with four nodes connected to two Nexus 3132Q-V cluster switches, CL1 and CL2.
- Cluster switch CL2 is to be replaced by C2 ([Steps 1 -21](#))
 - On each node, cluster LIFs connected to CL2 are migrated onto cluster ports connected to CL1.
 - Disconnect cabling from all ports on CL2 and reconnect cabling to the same ports on the replacement switch C2.
 - On each node, its migrated cluster LIFs are reverted.

Steps

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport

message:

```
system node autosupport invoke -node * -type all - message MAINT=xh
```

x is the duration of the maintenance window in hours.



The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Display information about the devices in your configuration:

```
network device-discovery show
```

```
cluster::> network device-discovery show
```

| Node | Local Port | Discovered Device | Interface | Platform |
|------|------------|-------------------|---------------|--------------|
| n1 | /cdp | | | |
| | e0a | CL1 | Ethernet1/1/1 | N3K-C3132Q-V |
| | e0b | CL2 | Ethernet1/1/1 | N3K-C3132Q-V |
| | e0c | CL2 | Ethernet1/1/2 | N3K-C3132Q-V |
| | e0d | CL1 | Ethernet1/1/2 | N3K-C3132Q-V |
| n2 | /cdp | | | |
| | e0a | CL1 | Ethernet1/1/3 | N3K-C3132Q-V |
| | e0b | CL2 | Ethernet1/1/3 | N3K-C3132Q-V |
| | e0c | CL2 | Ethernet1/1/4 | N3K-C3132Q-V |
| | e0d | CL1 | Ethernet1/1/4 | N3K-C3132Q-V |
| n3 | /cdp | | | |
| | e4a | CL1 | Ethernet1/7 | N3K-C3132Q-V |
| | e4e | CL2 | Ethernet1/7 | N3K-C3132Q-V |
| n4 | /cdp | | | |
| | e4a | CL1 | Ethernet1/8 | N3K-C3132Q-V |
| | e4e | CL2 | Ethernet1/8 | N3K-C3132Q-V |

12 entries were displayed

3. Determine the administrative or operational status for each cluster interface:

a. Display the network port attributes:

```
network port show
```

```
cluster::*> network port show -role cluster
(network port show)
```

```
Node: n1
```

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0c | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0d | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |

Node: n2

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0c | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0d | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |

Node: n3

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e4a | Cluster | Cluster | | up | 9000 | auto/40000 | - |
| - | | | | | | | |

```
e4e      Cluster      Cluster      up    9000 auto/40000 -
-

Node: n4

Ignore

Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
-----
e4a      Cluster      Cluster      up    9000 auto/40000 -
-
e4e      Cluster      Cluster      up    9000 auto/40000 -
-
12 entries were displayed.
```

b. Display information about the logical interfaces:

```
network interface show
```

```
cluster::*> network interface show -role cluster
(network interface show)
```

| | Current | Is | Logical | Status | Network | Current |
|---------|---------|------|-----------|------------|---------------|---------|
| Vserver | Port | Home | Interface | Admin/Oper | Address/Mask | Node |
| ----- | | | | | | |
| Cluster | | | | | | |
| e0a | | true | n1_clus1 | up/up | 10.10.0.1/24 | n1 |
| e0b | | true | n1_clus2 | up/up | 10.10.0.2/24 | n1 |
| e0c | | true | n1_clus3 | up/up | 10.10.0.3/24 | n1 |
| e0d | | true | n1_clus4 | up/up | 10.10.0.4/24 | n1 |
| e0a | | true | n2_clus1 | up/up | 10.10.0.5/24 | n2 |
| e0b | | true | n2_clus2 | up/up | 10.10.0.6/24 | n2 |
| e0c | | true | n2_clus3 | up/up | 10.10.0.7/24 | n2 |
| e0d | | true | n2_clus4 | up/up | 10.10.0.8/24 | n2 |
| e0a | | true | n3_clus1 | up/up | 10.10.0.9/24 | n3 |
| e0e | | true | n3_clus2 | up/up | 10.10.0.10/24 | n3 |
| e0a | | true | n4_clus1 | up/up | 10.10.0.11/24 | n4 |
| e0e | | true | n4_clus2 | up/up | 10.10.0.12/24 | n4 |

12 entries were displayed.

c. Display the information on the discovered cluster switches:

```
system cluster-switch show
```



```
cluster::> system cluster-switch show
```

| Switch | Type | Address | Model |
|--|-----------------|-------------|---------|
| CL1 | cluster-network | 10.10.1.101 | NX3132V |
| Serial Number: FOX000001 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |
| CL2 | cluster-network | 10.10.1.102 | NX3132V |
| Serial Number: FOX000002 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |

2 entries were displayed.

4. Verify that the appropriate RCF and image are installed on the new Nexus 3132Q-V switch as necessary for your requirements, and make any essential site customizations.

You must prepare the replacement switch at this time. If you need to upgrade the RCF and image, you must follow these steps:

- a. On the NetApp Support Site, go to the [Cisco Ethernet Switch](#) page.
 - b. Note your switch and the required software versions in the table on that page.
 - c. Download the appropriate version of the RCF.
 - d. Click **CONTINUE** on the **Description** page, accept the license agreement, and then follow the instructions on the **Download** page to download the RCF.
 - e. Download the appropriate version of the image software.
5. Migrate the LIFs associated to the cluster ports connected to switch C2:

```
network interface migrate
```

This example shows that the LIF migration is done on all the nodes:

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus2
-source-node n1 -destination-node n1 -destination-port e0a
cluster::*> network interface migrate -vserver Cluster -lif n1_clus3
-source-node n1 -destination-node n1 -destination-port e0d
cluster::*> network interface migrate -vserver Cluster -lif n2_clus2
-source-node n2 -destination-node n2 -destination-port e0a
cluster::*> network interface migrate -vserver Cluster -lif n2_clus3
-source-node n2 -destination-node n2 -destination-port e0d
cluster::*> network interface migrate -vserver Cluster -lif n3_clus2
-source-node n3 -destination-node n3 -destination-port e4a
cluster::*> network interface migrate -vserver Cluster -lif n4_clus2
-source-node n4 -destination-node n4 -destination-port e4a
```

6. Verify cluster's health:

```
network interface show
```

```
cluster::*> network interface show -role cluster
(network interface show)
```

| Current Is | Logical | Status | Network | Current | |
|------------|-----------|------------|---------------|---------|------|
| Vserver | Interface | Admin/Oper | Address/Mask | Node | Port |
| Home | | | | | |
| ----- | ----- | ----- | ----- | ----- | |
| ----- | ----- | | | | |
| Cluster | | | | | |
| | n1_clus1 | up/up | 10.10.0.1/24 | n1 | e0a |
| true | | | | | |
| | n1_clus2 | up/up | 10.10.0.2/24 | n1 | e0a |
| false | | | | | |
| | n1_clus3 | up/up | 10.10.0.3/24 | n1 | e0d |
| false | | | | | |
| | n1_clus4 | up/up | 10.10.0.4/24 | n1 | e0d |
| true | | | | | |
| | n2_clus1 | up/up | 10.10.0.5/24 | n2 | e0a |
| true | | | | | |
| | n2_clus2 | up/up | 10.10.0.6/24 | n2 | e0a |
| false | | | | | |
| | n2_clus3 | up/up | 10.10.0.7/24 | n2 | e0d |
| false | | | | | |
| | n2_clus4 | up/up | 10.10.0.8/24 | n2 | e0d |
| true | | | | | |
| | n3_clus1 | up/up | 10.10.0.9/24 | n3 | e4a |
| true | | | | | |
| | n3_clus2 | up/up | 10.10.0.10/24 | n3 | e4a |
| false | | | | | |
| | n4_clus1 | up/up | 10.10.0.11/24 | n4 | e4a |
| true | | | | | |
| | n4_clus2 | up/up | 10.10.0.12/24 | n4 | e4a |
| false | | | | | |

12 entries were displayed.

7. Shut down the cluster interconnect ports that are physically connected to switch CL2:

```
network port modify
```

This example shows the specified ports being shut down on all nodes:

```

cluster::*> network port modify -node n1 -port e0b -up-admin false
cluster::*> network port modify -node n1 -port e0c -up-admin false
cluster::*> network port modify -node n2 -port e0b -up-admin false
cluster::*> network port modify -node n2 -port e0c -up-admin false
cluster::*> network port modify -node n3 -port e4e -up-admin false
cluster::*> network port modify -node n4 -port e4e -up-admin false

```

8. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

```

cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e0a 10.10.0.1
Cluster n1_clus2 n1      e0b 10.10.0.2
Cluster n1_clus3 n1      e0c 10.10.0.3
Cluster n1_clus4 n1      e0d 10.10.0.4
Cluster n2_clus1 n2      e0a 10.10.0.5
Cluster n2_clus2 n2      e0b 10.10.0.6
Cluster n2_clus3 n2      e0c 10.10.0.7
Cluster n2_clus4 n2      e0d 10.10.0.8
Cluster n3_clus1 n4      e0a 10.10.0.9
Cluster n3_clus2 n3      e0e 10.10.0.10
Cluster n4_clus1 n4      e0a 10.10.0.11
Cluster n4_clus2 n4      e0e 10.10.0.12

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8 10.10.0.9 10.10.0.10
10.10.0.11 10.10.0.12
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 32 path(s)
Basic connectivity fails on 0 path(s)
.....
Detected 1500 byte MTU on 32 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.1 to Remote 10.10.0.9
    Local 10.10.0.1 to Remote 10.10.0.10
    Local 10.10.0.1 to Remote 10.10.0.11
    Local 10.10.0.1 to Remote 10.10.0.12

```

```
Local 10.10.0.2 to Remote 10.10.0.5
Local 10.10.0.2 to Remote 10.10.0.6
Local 10.10.0.2 to Remote 10.10.0.7
Local 10.10.0.2 to Remote 10.10.0.8
Local 10.10.0.2 to Remote 10.10.0.9
Local 10.10.0.2 to Remote 10.10.0.10
Local 10.10.0.2 to Remote 10.10.0.11
Local 10.10.0.2 to Remote 10.10.0.12
Local 10.10.0.3 to Remote 10.10.0.5
Local 10.10.0.3 to Remote 10.10.0.6
Local 10.10.0.3 to Remote 10.10.0.7
Local 10.10.0.3 to Remote 10.10.0.8
Local 10.10.0.3 to Remote 10.10.0.9
Local 10.10.0.3 to Remote 10.10.0.10
Local 10.10.0.3 to Remote 10.10.0.11
Local 10.10.0.3 to Remote 10.10.0.12
Local 10.10.0.4 to Remote 10.10.0.5
Local 10.10.0.4 to Remote 10.10.0.6
Local 10.10.0.4 to Remote 10.10.0.7
Local 10.10.0.4 to Remote 10.10.0.8
Local 10.10.0.4 to Remote 10.10.0.9
Local 10.10.0.4 to Remote 10.10.0.10
Local 10.10.0.4 to Remote 10.10.0.11
Local 10.10.0.4 to Remote 10.10.0.12
```

Larger than PMTU communication succeeds on 32 path(s)

RPC status:

8 paths up, 0 paths down (tcp check)

8 paths up, 0 paths down (udp check)

9. Shut down the ports 1/31 and 1/32 on CL1, and the active Nexus 3132Q-V switch:

shutdown

This example shows the ISL ports 1/31 and 1/32 being shut down on switch CL1:

```
(CL1)# configure
(CL1)(Config)# interface e1/31-32
(CL1(config-if-range)# shutdown
(CL1(config-if-range)# exit
(CL1)(Config)# exit
(CL1)#
```

10. Remove all the cables attached to the Nexus 3132Q-V switch CL2 and reconnect them to the replacement switch C2 on all nodes.

11. Remove the ISL cables from ports e1/31 and e1/32 on CL2 and reconnect them to the same ports on the replacement switch C2.
12. Bring up ISLs ports 1/31 and 1/32 on the Nexus 3132Q-V switch CL1.

```
(CL1)# configure
(CL1) (Config)# interface e1/31-32
(CL1(config-if-range)# no shutdown
(CL1(config-if-range)# exit
(CL1) (Config)# exit
(CL1)#
```

13. Verify that the ISLs are up on CL1:

```
show port-channel
```

Ports Eth1/31 and Eth1/32 should indicate (P) , which means that the ISL ports are up in the port-channel.

```
CL1# show port-channel summary
Flags: D - Down          P - Up in port-channel (members)
       I - Individual    H - Hot-standby (LACP only)
       s - Suspended     r - Module-removed
       S - Switched      R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met

-----
-----
Group Port-          Type  Protocol  Member                      Ports
Channel
-----
-----
1      Po1 (SU)       Eth    LACP     Eth1/31 (P)  Eth1/32 (P)
```

14. Verify that the ISLs are up on C2:

```
show port-channel summary
```

Ports Eth1/31 and Eth1/32 should indicate (P) , which means that both ISL ports are up in the port-channel.

```
C2# show port-channel summary
```

```
Flags: D - Down          P - Up in port-channel (members)
       I - Individual    H - Hot-standby (LACP only)
       s - Suspended     r - Module-removed
       S - Switched      R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
```

```
-----
-----
Group Port-          Type   Protocol  Member Ports
Channel
-----
-----
1      Po1 (SU)      Eth     LACP      Eth1/31 (P)  Eth1/32 (P)
```

15. On all nodes, bring up all the cluster interconnect ports connected to the Nexus 3132Q-V switch C2:
network port modify

```
cluster::*> network port modify -node n1 -port e0b -up-admin true
cluster::*> network port modify -node n1 -port e0c -up-admin true
cluster::*> network port modify -node n2 -port e0b -up-admin true
cluster::*> network port modify -node n2 -port e0c -up-admin true
cluster::*> network port modify -node n3 -port e4e -up-admin true
cluster::*> network port modify -node n4 -port e4e -up-admin true
```

16. For all nodes, revert all of the migrated cluster interconnect LIFs:

```
network interface revert
```

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus2
cluster::*> network interface revert -vserver Cluster -lif n1_clus3
cluster::*> network interface revert -vserver Cluster -lif n2_clus2
cluster::*> network interface revert -vserver Cluster -lif n2_clus3
Cluster::*> network interface revert -vserver Cluster -lif n3_clus2
Cluster::*> network interface revert -vserver Cluster -lif n4_clus2
```

17. Verify that the cluster interconnect ports are now reverted to their home:

```
network interface show
```

This example shows that all the LIFs are successfully reverted because the ports listed under the Current Port column have a status of true in the Is Home column. If the Is Home column value is false, the LIF has not been reverted.

```

cluster::*> network interface show -role cluster
(network interface show)

```

| Current Is | Logical | Status | Network | Current | |
|------------|-----------|------------|---------------|---------|-------|
| Vserver | Interface | Admin/Oper | Address/Mask | Node | Port |
| Home | | | | | |
| ----- | ----- | ----- | ----- | ----- | ----- |
| Cluster | | | | | |
| true | n1_clus1 | up/up | 10.10.0.1/24 | n1 | e0a |
| true | n1_clus2 | up/up | 10.10.0.2/24 | n1 | e0b |
| true | n1_clus3 | up/up | 10.10.0.3/24 | n1 | e0c |
| true | n1_clus4 | up/up | 10.10.0.4/24 | n1 | e0d |
| true | n2_clus1 | up/up | 10.10.0.5/24 | n2 | e0a |
| true | n2_clus2 | up/up | 10.10.0.6/24 | n2 | e0b |
| true | n2_clus3 | up/up | 10.10.0.7/24 | n2 | e0c |
| true | n2_clus4 | up/up | 10.10.0.8/24 | n2 | e0d |
| true | n3_clus1 | up/up | 10.10.0.9/24 | n3 | e4a |
| true | n3_clus2 | up/up | 10.10.0.10/24 | n3 | e4e |
| true | n4_clus1 | up/up | 10.10.0.11/24 | n4 | e4a |
| true | n4_clus2 | up/up | 10.10.0.12/24 | n4 | e4e |

12 entries were displayed.

18. Verify that the cluster ports are connected:

```
network port show
```

```

cluster::*> network port show -role cluster
(network port show)
Node: n1

Ignore

Speed (Mbps) Health

```



```

Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
-----
e0a       Cluster      Cluster      up    9000  auto/10000  -      -
e0b       Cluster      Cluster      up    9000  auto/10000  -      -
e0c       Cluster      Cluster      up    9000  auto/10000  -      -
e0d       Cluster      Cluster      up    9000  auto/10000  -      -

Node: n2

Ignore

Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
-----
e0a       Cluster      Cluster      up    9000  auto/10000  -      -
e0b       Cluster      Cluster      up    9000  auto/10000  -      -
e0c       Cluster      Cluster      up    9000  auto/10000  -      -
e0d       Cluster      Cluster      up    9000  auto/10000  -      -

Node: n3

Ignore

Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
-----
e4a       Cluster      Cluster      up    9000  auto/40000  -      -
e4e       Cluster      Cluster      up    9000  auto/40000  -      -

Node: n4

Ignore

Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
-----
-----
e4a       Cluster      Cluster      up    9000  auto/40000  -      -

```

```
e4e      Cluster      Cluster      up      9000 auto/40000 -      -  
12 entries were displayed.
```

19. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

```
cluster::*> cluster ping-cluster -node n1  
Host is n1  
Getting addresses from network interface table...  
Cluster n1_clus1 n1      e0a 10.10.0.1  
Cluster n1_clus2 n1      e0b 10.10.0.2  
Cluster n1_clus3 n1      e0c 10.10.0.3  
Cluster n1_clus4 n1      e0d 10.10.0.4  
Cluster n2_clus1 n2      e0a 10.10.0.5  
Cluster n2_clus2 n2      e0b 10.10.0.6  
Cluster n2_clus3 n2      e0c 10.10.0.7  
Cluster n2_clus4 n2      e0d 10.10.0.8  
Cluster n3_clus1 n3      e0a 10.10.0.9  
Cluster n3_clus2 n3      e0e 10.10.0.10  
Cluster n4_clus1 n4      e0a 10.10.0.11  
Cluster n4_clus2 n4      e0e 10.10.0.12  
  
Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4  
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8 10.10.0.9 10.10.0.10  
10.10.0.11 10.10.0.12  
Cluster Vserver Id = 4294967293  
Ping status:  
....  
Basic connectivity succeeds on 32 path(s)  
Basic connectivity fails on 0 path(s)  
.....  
Detected 1500 byte MTU on 32 path(s):  
    Local 10.10.0.1 to Remote 10.10.0.5  
    Local 10.10.0.1 to Remote 10.10.0.6  
    Local 10.10.0.1 to Remote 10.10.0.7  
    Local 10.10.0.1 to Remote 10.10.0.8  
    Local 10.10.0.1 to Remote 10.10.0.9  
    Local 10.10.0.1 to Remote 10.10.0.10  
    Local 10.10.0.1 to Remote 10.10.0.11  
    Local 10.10.0.1 to Remote 10.10.0.12  
    Local 10.10.0.2 to Remote 10.10.0.5  
    Local 10.10.0.2 to Remote 10.10.0.6  
    Local 10.10.0.2 to Remote 10.10.0.7  
    Local 10.10.0.2 to Remote 10.10.0.8  
    Local 10.10.0.2 to Remote 10.10.0.9
```

```
Local 10.10.0.2 to Remote 10.10.0.10
Local 10.10.0.2 to Remote 10.10.0.11
Local 10.10.0.2 to Remote 10.10.0.12
Local 10.10.0.3 to Remote 10.10.0.5
Local 10.10.0.3 to Remote 10.10.0.6
Local 10.10.0.3 to Remote 10.10.0.7
Local 10.10.0.3 to Remote 10.10.0.8
Local 10.10.0.3 to Remote 10.10.0.9
Local 10.10.0.3 to Remote 10.10.0.10
Local 10.10.0.3 to Remote 10.10.0.11
Local 10.10.0.3 to Remote 10.10.0.12
Local 10.10.0.4 to Remote 10.10.0.5
Local 10.10.0.4 to Remote 10.10.0.6
Local 10.10.0.4 to Remote 10.10.0.7
Local 10.10.0.4 to Remote 10.10.0.8
Local 10.10.0.4 to Remote 10.10.0.9
Local 10.10.0.4 to Remote 10.10.0.10
Local 10.10.0.4 to Remote 10.10.0.11
Local 10.10.0.4 to Remote 10.10.0.12
```

Larger than PMTU communication succeeds on 32 path(s)

RPC status:

8 paths up, 0 paths down (tcp check)

8 paths up, 0 paths down (udp check)

20. Display the information about the devices in your configuration:

- ° network device-discovery show
- ° network port show -role cluster
- ° network interface show -role cluster
- ° system cluster-switch show

```
cluster::> network device-discovery show
```

| Node | Local Port | Discovered Device | Interface | Platform |
|------|------------|-------------------|---------------|--------------|
| n1 | /cdp | | | |
| | e0a | C1 | Ethernet1/1/1 | N3K-C3132Q-V |
| | e0b | C2 | Ethernet1/1/1 | N3K-C3132Q-V |
| | e0c | C2 | Ethernet1/1/2 | N3K-C3132Q-V |
| | e0d | C1 | Ethernet1/1/2 | N3K-C3132Q-V |
| n2 | /cdp | | | |
| | e0a | C1 | Ethernet1/1/3 | N3K-C3132Q-V |
| | e0b | C2 | Ethernet1/1/3 | N3K-C3132Q-V |
| | e0c | C2 | Ethernet1/1/4 | N3K-C3132Q-V |
| | e0d | C1 | Ethernet1/1/4 | N3K-C3132Q-V |
| n3 | /cdp | | | |
| | e4a | C1 | Ethernet1/7 | N3K-C3132Q-V |
| | e4e | C2 | Ethernet1/7 | N3K-C3132Q-V |
| n4 | /cdp | | | |
| | e4a | C1 | Ethernet1/8 | N3K-C3132Q-V |
| | e4e | C2 | Ethernet1/8 | N3K-C3132Q-V |

```
12 entries were displayed.
```

```
cluster::*> network port show -role cluster
```

```
(network port show)
```

```
Node: n1
```

```
Ignore
```

| Health | Port | IPspace | Broadcast Domain | Link | MTU | Speed(Mbps) | Admin/Oper | Health |
|--------|------|---------|------------------|------|------|-------------|------------|--------|
| Status | | | | | | | | Status |
| | e0a | Cluster | Cluster | up | 9000 | auto/10000 | | - |
| | - | | | | | | | |
| | e0b | Cluster | Cluster | up | 9000 | auto/10000 | | - |
| | - | | | | | | | |
| | e0c | Cluster | Cluster | up | 9000 | auto/10000 | | - |
| | - | | | | | | | |
| | e0d | Cluster | Cluster | up | 9000 | auto/10000 | | - |
| | - | | | | | | | |

```
Node: n2
```

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e0a | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0b | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0c | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |
| e0d | Cluster | Cluster | | up | 9000 | auto/10000 | - |
| - | | | | | | | |

Node: n3

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e4a | Cluster | Cluster | | up | 9000 | auto/40000 | - |
| - | | | | | | | |
| e4e | Cluster | Cluster | | up | 9000 | auto/40000 | - |
| - | | | | | | | |

Node: n4

Ignore

| | | | | | | Speed(Mbps) | Health |
|--------|---------|-----------|--------|------|------|-------------|--------|
| Health | | | | | | | |
| Port | IPspace | Broadcast | Domain | Link | MTU | Admin/Oper | Status |
| Status | | | | | | | |
| ----- | ----- | ----- | | ---- | ---- | ----- | |
| ----- | ----- | | | | | | |
| e4a | Cluster | Cluster | | up | 9000 | auto/40000 | - |
| - | | | | | | | |
| e4e | Cluster | Cluster | | up | 9000 | auto/40000 | - |
| - | | | | | | | |

12 entries were displayed.

```
cluster::*> network interface show -role cluster
```

```
(network interface show)
```

| | | Logical | Status | Network | Current |
|------------|-----------|------------|--------------|---------------|---------|
| Current Is | | | | | |
| Vserver | Interface | Admin/Oper | Address/Mask | | Node |
| Port | Home | | | | |
| ----- | | | | | |
| ----- | | | | | |
| Cluster | | | | | |
| | | n1_clus1 | up/up | 10.10.0.1/24 | n1 |
| e0a | true | | | | |
| | | n1_clus2 | up/up | 10.10.0.2/24 | n1 |
| e0b | true | | | | |
| | | n1_clus3 | up/up | 10.10.0.3/24 | n1 |
| e0c | true | | | | |
| | | n1_clus4 | up/up | 10.10.0.4/24 | n1 |
| e0d | true | | | | |
| | | n2_clus1 | up/up | 10.10.0.5/24 | n2 |
| e0a | true | | | | |
| | | n2_clus2 | up/up | 10.10.0.6/24 | n2 |
| e0b | true | | | | |
| | | n2_clus3 | up/up | 10.10.0.7/24 | n2 |
| e0c | true | | | | |
| | | n2_clus4 | up/up | 10.10.0.8/24 | n2 |
| e0d | true | | | | |
| | | n3_clus1 | up/up | 10.10.0.9/24 | n3 |
| e4a | true | | | | |
| | | n3_clus2 | up/up | 10.10.0.10/24 | n3 |
| e4e | true | | | | |
| | | n4_clus1 | up/up | 10.10.0.11/24 | n4 |
| e4a | true | | | | |
| | | n4_clus2 | up/up | 10.10.0.12/24 | n4 |
| e4e | true | | | | |

```
12 entries were displayed.
```

```
cluster::*> system cluster-switch show
```

| Switch | Type | Address | Model |
|--|-----------------|-------------|---------|
| CL1 | cluster-network | 10.10.1.101 | NX3132V |
| Serial Number: FOX000001 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |
| CL2 | cluster-network | 10.10.1.102 | NX3132V |
| Serial Number: FOX000002 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |
| C2 | cluster-network | 10.10.1.103 | NX3132V |
| Serial Number: FOX000003 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |

3 entries were displayed.

21. Remove the replaced Nexus 3132Q-V switch, if it is not already removed automatically:

```
system cluster-switch delete
```

```
cluster::*> system cluster-switch delete -device CL2
```

22. Verify that the proper cluster switches are monitored:

```
system cluster-switch show
```

```
cluster::> system cluster-switch show
```

| Switch | Type | Address | Model |
|--|-----------------|-------------|---------|
| CL1 | cluster-network | 10.10.1.101 | NX3132V |
| Serial Number: FOX000001 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |
| C2 | cluster-network | 10.10.1.103 | NX3132V |
| Serial Number: FOX000002 | | | |
| Is Monitored: true | | | |
| Reason: | | | |
| Software Version: Cisco Nexus Operating System (NX-OS) Software, | | | |
| Version | | | |
| 7.0(3)I4(1) | | | |
| Version Source: CDP | | | |

2 entries were displayed.

23. Enable the cluster switch health monitor log collection feature for collecting switch-related log files:

```
system cluster-switch log setup-password
```

```
system cluster-switch log enable-collection
```



```

cluster::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
C1
C2

cluster::*> system cluster-switch log setup-password

Enter the switch name: C1
RSA key fingerprint is e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log setup-password

Enter the switch name: C2
RSA key fingerprint is 57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster::*>

```



If any of these commands return an error, contact NetApp support.

24. If you suppressed automatic case creation, re-enable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

Related information

[Cisco Ethernet Switch description page](#)

[Hardware Universe](#)

Copyright Information

Copyright © 2022 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.